

TL;DR

Across 1,067,371 rows of the UCI Online Retail II dataset, 25.6% of rows fail at least one High-severity or Blocker check; 22.8% have no Customer ID, and a currency column is absent entirely. I propose a 12-field intake specification — enforced in code via a pandera schema — plus six automated ingestion checks. The one improvement I would ship first is automated cancellation reconciliation + non-product-code filtering, which together account for 8.5% of nominal gross revenue and must be resolved before a pricing model can learn a truthful target.

1 · Data quality findings

Dataset: `online_retail_II.xlsx`, two sheets (`Year 2009-2010`, `Year 2010-2011`), 1,067,371 rows, period 2009-12-01 → 2011-12-09, 53,628 invoices, 5,304 stock codes, 5,942 customers, 43 countries.

I organised checks along six quality dimensions and framed every finding by its **pricing-model impact** rather than as an abstract defect. All numbers below are reproduced by `notebooks/02_quality_assessment.ipynb` (see `output/dq_results.csv`).


Severity-ranked findings

Dimension	Check	Failing rows	% of total	Severity	Pricing-model impact
Validity	Currency column absent	1,067,371	100.00%	Blocker	43 countries in the data; without an explicit currency column, GBP / EUR / USD are silently mixed as if comparable.
Completeness	Customer ID is null	243,007	22.77%	Blocker	Customer-level elasticity and retention segmentation both require a stable customer key. No safe imputation.
Uniqueness	Exact duplicate rows	12,133	1.14%	High	Duplicates inflate both revenue and volume, biasing aggregation-based features.
Validity	Price <= 0	6,207	0.58%	High	Non-positive unit prices contaminate the price distribution and bias learned elasticity.
Anomaly	Non-product stock codes (POST , DOT , BANK CHARGES , M , AMAZONFEE , TEST , ...)	5,810	0.54%	High	Fees and adjustments share the transactional schema but are not products; leaving them in dominates the training loss.
Integrity	Negative quantity on a non-cancellation invoice	3,457	0.32%	High	Outside the C - prefix convention, negative quantities are ambiguous — returns, corrections, and data errors become indistinguishable.
Integrity	Cancellation rows with no matching original purchase	2,234	0.21%	High	Unmatched cancellations either inflate returns or indicate lost origin records; either way,


Dimension	Check	Failing rows	% of total	Severity	Pricing-model impact
					net-revenue features are wrong.
Consistency	StockCodes with multiple distinct descriptions	420,566	39.40%	Medium	Description drift on the same SKU signals upstream master-data inconsistency and breaks text-based features.
Uniqueness	Duplicate (Invoice, StockCode) pairs	88,585	8.30%	Medium	Multiple lines of the same SKU on one invoice should be consolidated with summed quantity.
Consistency	StockCode casing / whitespace variants for the same product	66,980	6.28%	Medium	Same product under multiple SKU strings splits price-history and demand signal.
Completeness	Description is null or empty	8,764	0.82%	Medium	Blocks product clustering and human review of anomalies.
Anomaly	Price > 99.99th percentile (£2,555.82)	107	0.01%	Low	Manual-adjustment outliers at £10k+ skew log-transformed features.

Two guardrail checks (Price is null , Quantity == 0) returned zero failing rows and are kept in the suite for future handovers.

Visual summary



Missingness by column (1.07M rows)



Price distribution with the outlier tail separated at the 99.5th percentile.

What this blocks

A useful B2B pricing model needs three things the raw handover does not provide: (a) a stable **customer key** (missing for 22.8% of rows), (b) an explicit **currency** per transaction (missing entirely — 13.9% of rows are non-UK), and (c) a clean **product universe** with non-product codes stripped and cancellations reconciled. Of the gross nominal revenue of £19.29M across the two years, cancellations account for −£1.53M and non-product codes account for a further −£0.11M — together

≈8.5% of headline revenue that must be classified correctly before any elasticity, mix, or margin model is trained. The 6.3% SKU casing variants and 39.4% SKUs with drifting descriptions fragment the product dimension itself: without a normalisation layer, the same product appears as several distinct SKUs and its price history is silently split.

2 · Proposed intake specification

The spec is delivered in two forms. First, a one-page table a non-technical customer can fill in:

Field	Type	Required?	Format / Enum	Example	Automated check at ingestion
invoice_id	string	Yes	≤32 chars, uppercase / digits / hyphen	INV-0001	Regex + uniqueness across file
line_id	string	Yes	Unique with invoice_id	1	Composite-key uniqueness
sku_id	string	Yes	Uppercase alphanumeric, 2–32 chars	SKU-22423	Regex + normalisation audit
product_description	string	Yes	Non-empty	WHITE HANGING HEART T-LIGHT HOLDER	Non-empty + consistency per SKU
quantity	integer	Yes	Non-zero; negative only if is_cancellation=true	6	Range + cancellation-coherence check
unit_price	decimal	Yes	≥ 0	2.55	Range + outlier flag
currency	ISO-4217	Yes	3-letter code	GBP	Enum limited to customer's declared currencies
invoice_datetime	ISO-8601	Yes	UTC or with offset	2010-12-01T08:26:00Z	Range vs declared business period
customer_id	string	Conditional	Required for named customers	17850	Null permitted only for anonymous walk-in
country	ISO-3166-1 alpha-2	Yes	2-letter code	GB	Enum check
is_cancellation	boolean	Yes	Default false	false	Cross-field coherence with negative quantity
original_invoice_id	string	Conditional	Required if is_cancellation=true	INV-0001	Referential integrity to invoice_id

Second, the same spec is enforced in code as a `pandera DataFrameSchema` in `src/intake_schema.py`. The schema validates on ingestion and, on failure, returns a row-level report of violations — identical to what PRYZM's engineers currently produce by hand.

Six automated checks run at ingestion

1. **Schema validation** — types, nullability, enums (`pandera`).
2. **Range checks** — `unit_price ≥ 0`; `quantity ≠ 0`; `invoice_datetime` within the customer's declared business period.
3. **Referential integrity** — every row with `is_cancellation=true` resolves to an `original_invoice_id` that exists.
4. **Duplicate detection** — composite uniqueness on `(invoice_id, line_id)`.
5. **SKU normalisation audit** — report SKUs that differ only by casing or whitespace; fuzzy near-duplicate report for v2.
6. **Volume sanity bounds** — total row count and date span compared against what the customer declares in the onboarding form; a divergence > 5% is flagged for human review.

The notebook demonstrates both directions: a valid row passes, a raw-style row (`sku_id='85123a'`, `unit_price=-1.5`, `currency='gbp'`, `country='United Kingdom'`, `quantity=0`) fails with five distinct violations captured.

3 · Priority improvement: automated cancellation reconciliation + non-product filtering

The problem. Two classes of transaction contaminate the revenue signal a pricing model would train on: (i) cancellations — 19,494 rows prefixed with `C`, of which 2,234 have no matching original purchase for the same customer and SKU, and (ii) non-product stock codes — 5,810 rows using codes such as `POST`, `DOT`, `BANK CHARGES`, `M`, `AMAZONFEE`, `TEST`, `ADJUST`. Together they explain roughly £1.64M (≈8.5%) of nominal gross revenue. Leaving them in produces two different failure modes: orphan cancellations show up as unexplained demand drops, while fee lines show up as products whose “price” is wildly outside the distribution.

Why ship this first. Every other improvement — normalisation, missing-customer imputation, currency handling — compounds on top of a clean revenue baseline. This change is also the one with the clearest before/after metric for a non-technical stakeholder: *net revenue after reconciliation, $\Delta\%$ vs the raw file*.

v1 design (3–5 days, intern-scoped).

- **Pipeline position.** A new stage between raw ingestion and the pricing-model feature store, reading the validated intake dataframe and writing a `clean_transactions` table plus a `reconciliation_report` table.
- **Cancellation matching.** For each cancellation row, match on `(customer_id, sku_id, unit_price)` against the preceding 180 days of purchases for the same customer; apply FIFO offset to purchase quantities. Rows that cannot be matched are tagged `unmatched_cancellation` and surfaced in the report, not silently dropped.
- **Non-product filtering.** Maintain a versioned `non_product_codes.yaml` (seeded from the codes above) and route rows matching it to a `non_product_ledger` table instead of the product stream. Per-customer overrides supported.
- **Outputs.** Two tables and a one-page markdown report per customer: before/after gross, cancellation match rate, number of `non_product` rows, unmatched-cancellation sample, and any new non-product codes seen that are not yet in the YAML.
- **Success metric.** Match rate $\geq 90\%$ of cancellation rows; unmatched share $< 3\%$ of transactions; before/after gross delta reported and signed off by the onboarding lead.

v2 (later). Extend to fuzzy-matched cancellations (where price or SKU drifted between purchase and return), plug the reconciliation report into Grafana, and generalise the non-product YAML to a per-customer plug-in so pharma-specific adjustment codes (regulatory recalls, sample dispenses) can be handled without code changes.

Appendix · Reproducibility

Clone the repo, `pip install -r requirements.txt`, drop `online_retail_II.xlsx` into `data/`, and run `jupyter nbconvert --to notebook --execute notebooks/02_quality_assessment.ipynb` — every number in this document is emitted by that notebook. The intake schema is executable: `from src.intake_schema import validate`.